

## 1 Goal

The goal of the second assignment is to create a PostgreSQL data schema with 7 tables that are very similar to the tables that you created in Lab1. The tables have the same names, attributes, and data types as the tables of Lab1, and the same Primary Keys and Foreign Keys, but they also have some UNIQUE constraints and restrictions on NULL that are described below.

After you create the data schema with the 7 tables, you must write five SQL statements that use those tables. Under Resources→Lab2, we will provide you with data that you can load into your tables so that you can test the results of your queries. Testing can prove that a query is wrong, but not that it is right, so be careful. We will not give you the results of these queries on the load data; you should be able to figure out the results on that data yourselves. You can also test your queries on your own data. In the “real world”, you have to make and check your own tests.

## 2 Lab2 Schema

### 2.1 Create PostgreSQL Schema Lab2

You will create a Lab2 schema to set apart the database tables created in this lab from the tables you will create in future, as well as from tables (and other objects) in the default (public) schema. In PostgreSQL, a database can have more than one schema; see [here](#) for more details on PostgreSQL schemas. You create the Lab2 schema as follows:

```
CREATE SCHEMA Lab2;
```

Now that you have created the schema, you want to set Lab2 to be your default schema when you use psql. If you do not set Lab2 as the default schema, then you will have to qualify your table names with the schema name (e.g., Lab2.Passenger). To set the default schema, you modify your search path. (For more details, see [here](#).)

```
ALTER ROLE cse180 SET SEARCH_PATH to Lab2;
```

You will need to log out and log back in to the server for this default schema change to take effect.  
(Students often forget to do this.)

You do not have to include the CREATE SCHEMA or ALTER ROLE statements in your solution.

### 2.2 Create tables

As in Lab1, you will create the tables **Passenger**, **Airport**, **Flight**, **CrewMember**, **Reservation**, **CheckIn**, and **FlightCrewAssignment**. The attributes for these tables are the same as for the tables of Lab1. Moreover, the data types for the attributes in these tables are the same as the ones specified for the tables of Lab1, and the Primary Keys and other Foreign Keys are also the same. You may use the *create\_lab1.sql* file solution that we provided on Piazza for Lab2 as the basis for the *create\_lab2.sql* file which you include in your Lab2 solution (although you don't have to include our comments, and you're using Lab2 instead of Lab1 as the schema). However, your Lab2 tables must have additional constraints, which are described in the next section.

### 2.2.1 Constraints

The following attributes cannot be NULL. All other attributes can be NULL ... but remember that attributes in Primary Keys also cannot be NULL, even though NOT NULL isn't specified for them.

- In Passenger: passengerName
- In Airport: city
- In Flight: actualDeparture, actualArrival
- In CrewMember: crewName

Also, the following must be unique for the specified table. That is, there cannot be identical rows in that table that have exactly the same (non-NULL) values for all of those attributes (composite unique constraint).

- In Passenger: the 2 attributes passengerName and email
- In Airport: the 2 attributes city and country
- In CrewMember: the attribute crewName
- In Reservation: the 3 attributes bookingDate, paymentMethod, and passengerID

Explanations:

- The first constraint says that there can't be two rows in Passenger which have the same values for both of the attributes passengerName and email (if both of those attributes are not NULL).

For example, there can't be two passengers whose passengerName is 'John Doe', and whose email is 'john@example.com'. But there could be two passengers whose passengerName is 'John Doe' and whose email is NULL.

- The second constraint says that there can't be two rows in Airport which have the same values for city and country (recall that city here isn't allowed to be NULL).
- The third constraint says that there can't be two rows in CrewMember which have the same values for crewName (recall that crewName here isn't allowed to be NULL).
- The fourth constraint says that there can't be two rows in Reservation which have the same values for all three of the attributes bookingDate, paymentMethod, and passengerID (if all three of those attributes are not NULL).

This means that we disallow duplicate "same-day, same payer, same method" bookings that appear to be the same transaction.

You will write a CREATE TABLE command for each of the 7 tables which has these additional constraints. Save the commands in the file *create\_lab2.sql*

### 3 SQL Queries

Below are English descriptions of the five SQL queries that you need to write for this assignment, which you will include in files queryX.sql, where X is the number of the query, e.g., your SQL statement for Query 1 will be in the file query1.sql, and so forth. Follow the directions as given. You will lose points if you give extra tuples or attributes in your results, if you give attributes with the wrong names or in the wrong order, or if you have missing or wrong results. You will also lose points if your queries are unnecessarily complex, even if they are correct. Grading is based on correctness of queries on all data, not just the load data that we have provided.

Remember that the Referential Integrity (foreign key) constraints from Lab1 are still in effect here. For example, any reservationID that is in a CheckIn row must appear as a reservationID in the Reservation table.

Attributes should have their original names in the results of your queries, unless an alias is requested. And if a query asks that several attributes appear in the result, the first attribute mentioned should appear first, the second attribute mentioned should appear second, etc.

#### 3.1 Query 1

Write a SQL query which finds the reservations for which the name of the passenger starts with 'R' (with that capitalization), the departure airport has exactly 12 runways, and the ticket price paid is greater than 40.00 dollars. The attributes in your result should be the name of the passenger, the (departure) airport code, and the ticket price, which should appear in your result as thePassengerName, theAirportCode, and theTicketPrice.

No duplicates should appear in your result.

#### 3.2 Query 2

A reservation is paid with a Visa card if the payment method (paymentMethod) is 'V' (uppercase V).

Write a SQL query which finds all passengers who have made a reservation paid with a Visa and who have never (in any flight and using any payment method) made a reservation that exceeds 100.00 dollars in ticket price. The attributes in your result should be passengerID and date of birth, appearing as thePassengerID and theDateOfBirth. Tuples in your result should be in reverse order by theDateOfBirth; if two tuples have the same date of birth, they should appear in increasing order of thePassengerID.

No duplicates should appear in your result.

**3.3 Query 3**

Recall that an assignment of a crew member to a flight is recorded in FlightCrewAssignment.

Write a SQL query which finds all crew members, excluding 'Sophia Kim', who were assigned to the exact same aircraft type on at least 2 different flights. The attribute in your result should be the crew member's ID, which should appear as the CrewID.

No duplicates should appear in your result.

(Mapping note: exclude rows where CrewMember.crewName = 'Sophia Kim'.)

**3.4 Query 4**

Note: If myTimestamp is a TIMESTAMP, then in PostgreSQL, DATE(myTimestamp) is the DATE value which is in that timestamp. Unfortunately, different relational database systems have different ways of extracting the DATE from a TIMESTAMP. (There are at least 3 ways to do that in PostgreSQL; this one is the simplest.)

Recall that a tuple in CheckIn identifies a passenger checking in for a flight with an associated reservation. Write a SQL query which finds all reservations for which the passenger checked-in and for which all of the following are true:

- The seat class requires special handling (seatClass = 'F').
- The date of checkInTime is January 12, 2024 or later.
- The departure airport location has 'en' appearing anywhere in its city, with that capitalization.
- The bagCount is greater than or equal to 2.
- The payment method used for the associated reservation is NULL.

The attributes in your result should be the flight ID and the check-in time, which should appear as theFlightID and theCheckInTime.

No duplicates should appear in your result.

**3.5 Query 5**

Several flight assignments in the FlightCrewAssignment table might be linked to flights that share the same actual departure date. We want to find all the assignments which are associated with the earliest actual departure date. An assignment has the earliest date if there are no assignments whose associated flight has an earlier actual departure date.

Write a SQL query which finds the flightID and crewID for all assignments whose associated flight has the earliest actual departure date. The attributes in your result should appear as theFlightID and theCrewID.

No duplicates should appear in your result.

#### 4 Testing

While your solution is still a work in progress, it is a good idea to drop all objects from the database every time you run the script, so you can start fresh. Of course, dropping each object may be tedious, and sometimes there may be a particular order in which objects must be dropped. The following commands (which you can put at the top of *create\_lab2.sql* if you want, but you don't have to), will drop your Lab2 schema (and all objects within it), and then create the (empty) schema again:

```
DROP SCHEMA Lab2 CASCADE;  
CREATE SCHEMA Lab2;
```

Before you submit, login to your database via psql and execute your script. As you've learned already, the command to execute a script is: \i <filename>.

Under Resources→Lab2 on Piazza, we will post a load script named *load\_lab2.sql* that loads data into the 7 tables of the database. You can execute that script with the command:

```
\i load_lab2.sql
```

You can test your 5 queries using that data, but you will have to figure out on your own whether your query results are correct. We won't provide answers to the 5 queries when they are run on the load data, and students should not share answers with other students. Also, your queries must be correct on any database instance, not just on the data that we provide. You may want to test your SQL statements on your own data as well.

## 5 Submitting

1. Save your scripts for table creations and query statements as *create\_lab2.sql* and *query1.sql* through *query5.sql*. You may add informative comments inside your scripts if you want (lines that start with two hyphens are interpreted as comment lines).
2. Zip the file(s) to a single file with name Lab2\_XXXXXXX.zip where XXXXXXX is your 7-digit student ID. For example, if a student's ID is 1234567, then the file that this student submits for Lab2 should be named Lab2\_1234567.zip

To generate the zip file you can use the Unix command:

```
zip Lab2_1234567 create_lab2.sql query1.sql query2.sql query3.sql query4.sql query5.sql
```

(Of course, you use your own student ID, not 1234567.)

3. Lab2 is due by 11:59pm on Sunday, October 26. Late submissions will not be accepted, and there will be no make-up Lab assignments.
4. You can get always rid of duplicates by using DISTINCT in your SELECT. But you will **lose points** if you use DISTINCT in cases where it is not necessary (i.e., if the query cannot return duplicates). You will also lose points if you weren't told to eliminate duplicates but you do eliminate them.
5. Be sure to follow directions about Academic Integrity that are in the Syllabus. If you have any questions about those directions, please speak to the instructor as soon as possible.